

CSCI-UA.60-1

Database Design

and Implementation

An Introduction to MongoDB

Prof Deena Engel

Department of Computer Science

deena.engel@nyu.edu

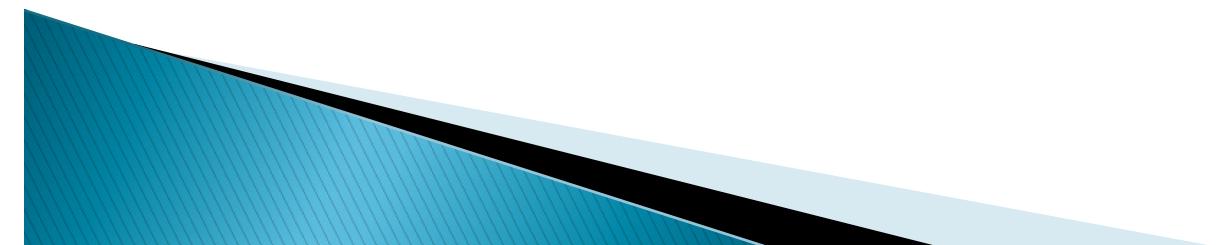
Getting started!

- ▶ Principles and concepts to remember with MongoDB:
 - A **document** is the basic unit of data
 - A **collection** can be thought of ... as a dynamic table
 - One instance of MongoDB can host multiple independent **databases** (and then each database can hold its own collections)
 - Every document has a special key (**_id**) which is unique within its collection
 - MongoDB comes with a simple but powerful **JavaScript shell** (but you do not need to know JavaScript for the work we will do in this class).



Getting started!

- ▶ MongoDB comes with a simple but powerful **JavaScript shell** (but you do not need to know JavaScript for the work we will do in this class).
- ▶ See
<https://docs.mongodb.com/manual/reference/method/> for built-in JavaScript methods.



“Hello, world!”

```
{ “greeting” : “Hello, world!” }
```

This simple document contains a single key (“**greeting**”) with the value of “**Hello, world!**”.



Rules about keys

- ▶ Keys must not contain the \0 (or null) character (because that is used to signify the end of a key).
- ▶ The period (.) and \$ characters have special properties and should be considered *reserved*.
- ▶ Note that a single document may not contain duplicate keys.



Collections

- ▶ Collections have *dynamic schemas* ...
- ▶ However, in reality, it is important to keep different kinds of documents in different collections.
- ▶ This will allow for ...
 - Faster searches
 - ... and other considerations.



Collection Names

- ▶ Collection names can use any valid UTF-8 string.
 - UTF-8 can represent any character in the Unicode standard and is backwardsly compatible with ASCII
 - http://www.w3schools.com/tags/ref_charactersets.asp
- ▶ Collection names may not contain the null character (\0) and should not use \$ in the name



Databases

- ▶ Collections are grouped into databases.
- ▶ Unlike with MySQL, you may run one database through i6.cims.nyu.edu within your account.
- ▶ Your account name and database name are the same.



Some commands to get you started ...

- ▶ Once you have opened your mongoDB session:

- ▶ To enter your database:

```
use <database name>
```

- ▶ Note: The following does not work for us, as everyone has one database:

```
show dbs
```

Class notes

- ▶ As with our other databases, I will post both class transcripts (records of the sessions) and the data files that we use to the course home pages under the “Class Notes” tab.



Getting help!

- ▶ To get documentation on the methods you can use with your database:

`db.help()`

- ▶ To get documentation on the methods you can use with your collection:

`db.users.help()`

- ▶ For example:

> `db.users.help()`



Working with external data files

- ▶ We will import data using examples with both **.json** and **.csv** files
- ▶ For example, to import a **.json** file called **users.json** into a collection called **users** within a database called **test** through a port numbered **1815**:

```
mongoimport --db your_database_name --collection  
your_collection_name --host class-  
mongodb.cims.nyu.edu --username your_username --  
password your_password --type json --file  
your_json_file_name
```

find() method

- ▶ The find() method takes two arguments:
 1. criteria
 2. projection
- ▶ Some examples are:
 - `db.books.find()`
 - - *no criteria, no selection!*
 - `db.books.find({ "AUTHOR": "Tolstoy", "Leo" })`
 - will find books by Leo Tolstoy
 - `db.books.find({ "PRICE": { $gt: 25 } })`
 - will find books over \$25

count()

- ▶ Note that these two statements are equivalent:
 - `db.books.find().count()`
 - `db.books.count()`
- ▶ To find all of the books published by Penguin in our collection:
 - `db.books.find({ "EDITION": "Penguin" }).count()`
- ▶ You can also include the criteria in *count()*:
 - `db.books.count({ "EDITION": "Penguin" })`



mongoDB Query operators

- ▶ \$gt greater than query.
- ▶ \$gte equal to or greater than
- ▶ \$in Matches any of the values that exist in an array specified in the query.
- ▶ \$lt less than
- ▶ \$lte less than or equal to
- ▶ \$ne not equal query.
- ▶ \$nin Matches values that are not in an array specified to the query.

Using \$or

- ▶ The mongo \$or takes an array of conditions as its argument.
 - <http://docs.mongodb.org/manual/reference/operator/query/or/>

\$in and \$nin

- ▶ To use **\$in**, create an array and then test the values of the data:
 - `db.books.find({ "YEAR_WRITTEN": { $in: [1865, 1925] } })`
- ▶ **\$nin** works in the same way, but to exclude with “not in”.

MongoDB's Aggregation Framework

- ▶ MongoDB offers a powerful aggregation framework.
- ▶ These queries are somewhat similar to the **GROUP BY** statements that we used in MySQL.
- ▶ The aggregation framework is only available in versions 2.4+
- ▶ See
<http://docs.mongodb.org/manual/aggregation/>

Class Notes

- ▶ Examples will be included in class discussion along with the data files so that you can practice as well.
- ▶ Some of the aggregation framework `$group` operators are:
 - `$min`
 - `$max`
 - `$sum`
 - `$avg`
- ▶ There are many more; see reference at <http://docs.mongodb.org/manual/reference/operator/aggregation/>



GeoSpatial Data in MongoDB

- ▶ MongoDB natively supports geo-spatial data (e.g. latitude & longitude) and offers many ways to work with these data.
- ▶ This is an example that corresponds to our discussion in class about a case where MongoDB requires a built-in index to work with the data.
- ▶ In our brief example, we will use the “2d” index. MongoDB offers a “2dsphere” index as well.
- ▶ For example, here is documentation on `$near`:
<http://docs.mongodb.org/manual/reference/operator/query/near/>
- ▶ The data under discussion in class are available on the course Readings pages.

Sample queries using geo-spatial data

- ▶ Latitude and longitude are stored in an array of 2 elements. For example, the Courant Institute location would be represented as follows:

[40.7296,-73.995305]

- ▶ Assuming the location of a place in our “wifi” collection is represented by the values associated with the key “Shape” in a document, one could create the following index:

```
db.wifi.ensureIndex( { "Shape":"2d"})
```



Finding the wifi locations closest to the Courant Institute!

- ▶ Using the \$near operator:

```
db.wifi.find( { "Shape" :  
    { $near : [ 40.7296,-73.995305 ] } } )
```

- ▶ Using the \$near operator and only displaying the name and address of the resulting data:

```
db.wifi.find( { "Shape" : { $near : [ 40.7296,-73.995305 ] } },  
    { _id:0,"NAME":1,"ADDRESS":1} )
```

- ▶ Using the \$near operator and only displaying the name and address of the three closest wifi hotspots to the Courant Institute:

```
db.wifi.find( { "Shape" : { $near : [ 40.7296,-73.995305 ] } },  
    { _id:0,"NAME":1,"ADDRESS":1} ).limit(3)
```

- ▶ We will look at more examples in class!